

HarvardX PH125.9 - Data Science: Capstone Project - Movielens

Reema Sajnani

2024-05-11

Contents

1	EXECUTIVE SUMMARY	2
2	STARTING CODE	3
2.1	Load Required Packages & Libraries	3
3	DATA EXPLORATION & VISUALIZATION	6
3.1	Data Class	6
3.2	Missing Values	6
3.3	Movies' Ratings	7
3.4	Users' Ratings	8
3.5	Genre Ratings	10
4	DATA MODELLING	12
4.1	Basic Model	12
4.2	Movie Effect Model	12
4.3	Movie & User Effect Model	13
4.4	Regularized Movie & User Effect Model	13
5	RESULTS & CONCLUSION	15
6	LIMITATIONS	15

1 EXECUTIVE SUMMARY

This project aims to produce a movie recommendation system, based on the 10M version of the *Movielens* data set from the *dslabs* package. The data set contains ratings of movies given by individuals over time. Using the code provided in the **HarvardX PH125.9 - Data Science: Capstone** course as a starting point, the data set is split into the **edx** partition (used to create the prediction algorithm) and the **final_holdout_test** partition (used to evaluate the aforementioned algorithm results). We first explore the data and analyze each input to notice some trends and patterns. We then use the inputs to evaluate their influence on our predictions. We then put our prediction models to test and choose the one that minimizes the RMSE (Root Mean Squared Error).

$$RMSE = \sqrt{\frac{1}{N} \sum_{u,i} (\hat{y}_{u,i} - y_{u,i})^2}$$

The best resulting code achieved in the project delivers an RMSE of **0.8641**, which was validated against a target of **0.8649**.

2 STARTING CODE

2.1 Load Required Packages & Libraries

Note: this process could take a couple of minutes

```
if(!require(tidyverse)) install.packages("tidyverse", repos = "http://cran.us.r-project.org")

## Loading required package: tidyverse

## Warning: package 'tidyverse' was built under R version 4.3.3

## Warning: package 'ggplot2' was built under R version 4.3.3

## Warning: package 'dplyr' was built under R version 4.3.3

## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr      1.1.4      v readr      2.1.4
## v forcats    1.0.0      v stringr    1.5.0
## v ggplot2     3.5.1      v tibble     3.2.1
## v lubridate  1.9.2      v tidyr      1.3.0
## v purrr      1.0.2

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors

if(!require(caret)) install.packages("caret", repos = "http://cran.us.r-project.org")

## Loading required package: caret

## Warning: package 'caret' was built under R version 4.3.3

## Loading required package: lattice

## Warning: package 'lattice' was built under R version 4.3.3

##
## Attaching package: 'caret'
##
## The following object is masked from 'package:purrr':
##
##   lift

if(!require(caret)) install.packages("knitr", repos = "http://cran.us.r-project.org")

library(tidyverse)

library(caret)

library(knitr)
```

```
## Warning: package 'knitr' was built under R version 4.3.3
```

```
#MovieLens 10M dataset:
```

```
#https://grouplens.org/datasets/movielens/10m/
```

```
#http://files.grouplens.org/datasets/movielens/ml-10m.zip
```

```
options(timeout = 120)
```

```
dl <- "ml-10M100K.zip"
```

```
if(!file.exists(dl))
```

```
  download.file("https://files.grouplens.org/datasets/movielens/ml-10m.zip", dl)
```

```
ratings_file <- "ml-10M100K/ratings.dat"
```

```
if(!file.exists(ratings_file))
```

```
  unzip(dl, ratings_file)
```

```
movies_file <- "ml-10M100K/movies.dat"
```

```
if(!file.exists(movies_file))
```

```
  unzip(dl, movies_file)
```

```
ratings <- as.data.frame(str_split(read_lines(ratings_file), fixed("::"), simplify = TRUE),  
                        stringsAsFactors = FALSE)
```

```
colnames(ratings) <- c("userId", "movieId", "rating", "timestamp")
```

```
ratings <- ratings %>%
```

```
  mutate(userId = as.integer(userId),  
         movieId = as.integer(movieId),  
         rating = as.numeric(rating),  
         timestamp = as.integer(timestamp))
```

```
movies <- as.data.frame(str_split(read_lines(movies_file), fixed("::"), simplify = TRUE),  
                        stringsAsFactors = FALSE)
```

```
colnames(movies) <- c("movieId", "title", "genres")
```

```
movies <- movies %>%
```

```
  mutate(movieId = as.integer(movieId))
```

```
movielens <- left_join(ratings, movies, by = "movieId")
```

```
# Final hold-out test set will be 10% of MovieLens data
```

```
set.seed(123, sample.kind="Rounding")
```

```
## Warning in set.seed(123, sample.kind = "Rounding"): non-uniform 'Rounding'
```

```
## sampler used
```

```
test_index <- createDataPartition(y = movielens$rating, times = 1, p = 0.1, list = FALSE)
```

```
edx <- movielens[-test_index,]
```

```
temp <- movielens[test_index,]
```

```
# Make sure userId and movieId in final hold-out test set are also in edx set
final_holdout_test <- temp %>% semi_join(edx, by = "movieId") %>% semi_join(edx, by = "userId")

# Add rows removed from final hold-out test set back into edx set
removed <- anti_join(temp, final_holdout_test)

## Joining with 'by = join_by(userId, movieId, rating, timestamp, title, genres)'

edx <- rbind(edx, removed)

rm(dl, ratings, movies, test_index, temp, movielens, removed)
```

3 DATA EXPLORATION & VISUALIZATION

3.1 Data Class

After the data is split into the training set, named **edx**, the first step is to understand the data and each of the 6 vectors.

- **userId** <integer> unique identifier per person
- **movieId** <numeric> unique identifier per movie
- **rating** <numeric> a rating on the scale of 0 to 5 per movie per user
- **timestamp** <integer> exact time of rating
- **title** <character> name and year of movie release
- **genres** <character> genre tags for each movie

Each of the nearly 9 million rows captures a rating given by one of the 69,878 unique users to one of the 10,677 unique movies in the library.

```
##      userId movieId rating timestamp                title
## 1         1     122      5 838985046             Boomerang (1992)
## 3         1     231      5 838983392             Dumb & Dumber (1994)
## 5         1     316      5 838983392             Stargate (1994)
## 6         1     329      5 838983392 Star Trek: Generations (1994)
## 7         1     355      5 838984474             Flintstones, The (1994)
## 8         1     356      5 838983653             Forrest Gump (1994)
##                                     genres
## 1                               Comedy|Romance
## 3                               Comedy
## 5             Action|Adventure|Sci-Fi
## 6 Action|Adventure|Drama|Sci-Fi
## 7             Children|Comedy|Fantasy
## 8             Comedy|Drama|Romance|War

##      n_users n_movies
## 1    69878    10677
```

3.2 Missing Values

A summary view of the data set also confirms no missing values.

```
##      userId      movieId      rating      timestamp
## Min.   :    1  Min.   :    1  Min.   :0.500  Min.   :7.897e+08
## 1st Qu.:18127  1st Qu.:   648  1st Qu.:3.000  1st Qu.:9.468e+08
## Median :35750  Median :  1834  Median :4.000  Median :1.035e+09
## Mean   :35873  Mean   :   4122  Mean   :3.512  Mean   :1.033e+09
## 3rd Qu.:53609  3rd Qu.:  3624  3rd Qu.:4.000  3rd Qu.:1.127e+09
## Max.   :71567  Max.   :65133  Max.   :5.000  Max.   :1.231e+09
##      title      genres
## Length:9000062  Length:9000062
## Class :character  Class :character
## Mode  :character  Mode  :character
##
##
##
```


3.3 Movies' Ratings

The most reviewed movies seem to be popular movie names from the 90s, hinting that re-known films are probably likely to be rated more often than lesser heard of movies.

```
## 'summarise()' has grouped output by 'movieId'. You can override using the
## '.groups' argument.
```

```
## # A tibble: 10,677 x 3
## # Groups:   movieId [10,677]
##   movieId title                                     count
##   <int> <chr>                                     <int>
## 1     296 Pulp Fiction (1994)                   31471
## 2     356 Forrest Gump (1994)                   30948
## 3     593 Silence of the Lambs, The (1991)       30336
## 4     480 Jurassic Park (1993)                   29332
## 5     318 Shawshank Redemption, The (1994)       27982
## 6     110 Braveheart (1995)                     26140
## 7     457 Fugitive, The (1993)                   26105
## 8     589 Terminator 2: Judgment Day (1991)      26079
## 9     260 Star Wars: Episode IV - A New Hope (a.k.a. Star Wars) (1977) 25715
## 10    150 Apollo 13 (1995)                       24347
## # i 10,667 more rows
```

The data also shows that lower ratings are less common and so are half star ratings. The most common ratings are 4,3 and 5 respectively



On closer examination, we will also notice that some movies are not rated as often as others, in fact, 126 movies are rated only once. More so, the least frequently rated movies titles and their average ratings appear to be obscure. This calls for regularization and a penalty to be applied in the calculation of the model.

Number of Ratings Per Movie





Table 1: Least Frequently Rated Movies

title	count	rating
100 Feet (2008)	1	2.0
4 (2005)	1	2.5
Accused (Anklaget) (2005)	1	0.5
Ace of Hearts (2008)	1	2.0
Ace of Hearts, The (1921)	1	3.5
Adios, Sabata (Indio Black, sai che ti dico: Sei un gran figlio di...) (1971)	1	1.5
Africa addio (1966)	1	3.0
Bad Blood (Mauvais sang) (1986)	1	4.5
Battle of Russia, The (Why We Fight, 5) (1943)	1	3.5
Besotted (2001)	1	0.5


3.4 Users' Ratings

The graph showing how often users rate movies is skewed to the right, implying that some users rarely rate movies, while some others rate thousands of movies.



We can also infer that some users are more critical in their ratings and some of them might have not rated too many movies. This phenomenon, as well, calls for a penalty to be applied in the calculation of the model.

Histogram of Average Movie Ratings By User



3.5 Genre Ratings

In the data set, we have 797 unique combinations of genres. When looking into the highest rated genre combinations, we see that most of the top rated ones are also the more frequently rated ones. Genre combinations with lower average ratings seem to be rated less often.

Table 2: Genres that Have Received the Highest Average Ratings

genres	count	rating
Animation IMAX Sci-Fi	7	4.714286
Adventure Fantasy Film-Noir Mystery Sci-Fi	3	4.333333
Drama Film-Noir Romance	3040	4.313651
Action Crime Drama IMAX	2358	4.292197
Animation Children Comedy Crime	7122	4.279837
Film-Noir Mystery	6014	4.235367
Crime Film-Noir Mystery	4010	4.227681
Film-Noir Romance Thriller	2431	4.218840
Crime Film-Noir Thriller	4860	4.207407
Crime Mystery Thriller	26945	4.199146

Table 3: Genres that Have Received the Lowest Average Ratings

genres	count	rating
Documentary Horror	638	1.464733
Action Animation Comedy Horror	2	1.500000
Action Horror Mystery Thriller	320	1.581250
Action Drama Horror Sci-Fi	5	1.600000
Comedy Film-Noir Thriller	24	1.729167
Adventure Drama Horror Sci-Fi Thriller	220	1.777273
Action Horror Mystery Sci-Fi	19	1.868421
Action Adventure Drama Fantasy Sci-Fi	58	1.879310
Action Children Comedy	524	1.887405
Children Fantasy Sci-Fi	57	1.903509

4 DATA MODELLING

The model relies on the statistical concept of RMSE (Root Mean Squared Error) and Regularization.

$$RMSE = \sqrt{\frac{1}{N} \sum_{u,i} (\hat{y}_{u,i} - y_{u,i})^2}$$

RMSE represents the error made while predicting a movie rating. Our goal is to model an algorithm such that RMSE is minimized as much as possible, while being under the goal of **0.8649**.

4.1 Basic Model

This is the simplest approach that predicts the same rating for all movies and is based on the average of ratings in the training data set.

The formula used for this model is:

$$Y_{u,i} = \hat{\mu} + \varepsilon_{u,i}$$

where $\hat{\mu}$ is the mean of all ratings (3.5125) and $\varepsilon_{i,u}$ is the random errors sampled from the same 0 centered distribution.

```
##           Model      RMSE
## 1 Basic Model 1.060057
```

The RMSE on the `final_holdout_test` data set is **1.06**. It is quite distant from the target of 0.8649 and indicates that this model doesn't produce ideal results.

4.2 Movie Effect Model

From exploration of the `edx` data set previously, we understood that some movies are not rated as often. We also noticed some obscurities in the least frequently rated movies.

The formula used for this model is:

$$Y_{u,i} = \hat{\mu} + b_i + \epsilon_{u,i}$$

where $\hat{\mu}$ is the mean of all ratings (3.5125) and $\varepsilon_{i,u}$ is the random errors sampled from the same 0 centered distribution. The b_i is a measure of the degree of popularity bias of each movie i .

```
##           Model      RMSE
## 1      Basic Model 1.0600565
## 2 Movie Effect Model 0.9431062
```

We used the impact of the movie ratings to calculate the RMSE of 0.9431, which is still much higher than the target of 0.8649

4.3 Movie & User Effect Model

In addition to the movie effect, we had also noticed an impact from users. Some users were rating less frequently and were also inclined to give movies lower ratings.

The formula used for this model is:

$$Y_{u,i} = \hat{\mu} + b_i + b_u + \epsilon_{u,i}$$


where $\hat{\mu}$ is the mean of all ratings (3.5125) and $\epsilon_{i,u}$ is the random errors sampled from the same 0 centered distribution. The b_i is a measure of the degree of popularity bias of each movie i . The b_u is a measure of the degree of user bias of each movie u .

##	Model	RMSE
## 1	Basic Model	1.0600565
## 2	Movie Effect Model	0.9431062
## 3	Movie + User Effect Model	0.8645388

We used the impact of the movie and user ratings to calculate the RMSE of 0.8645, which is lower than the target of 0.8649, hence proving that this model is quite good. However, there might be room to better here.

4.4 Regularized Movie & User Effect Model

From the Movie and User Effects Models we have seen the influence these inputs have on our predictions. Therefore it becomes necessary to tune our model, calling for the use of Regularization. Essentially, regularization involves incorporating a penalty for high values of b_i and b_u into the sum of squares equation that we aim to minimize. The penalization or tuning parameter is called *Lambda*



The optimal Lambda turns out to be 5.25. At this value of Lambda, RMSE is minimized to **0.8641** which is even lower than our last model.

##	Model	RMSE
## 1	Basic Model	1.0600565
## 2	Movie Effect Model	0.9431062
## 3	Movie + User Effect Model	0.8645388
## 4	Regularized Movie + User Based Model	0.8641476

5 RESULTS & CONCLUSION

The RMSE values of all the represented models are the following:

##	Model	RMSE
## 1	Basic Model	1.0600565
## 2	Movie Effect Model	0.9431062
## 3	Movie + User Effect Model	0.8645388
## 4	Regularized Movie + User Based Model	0.8641476

We developed a machine learning algorithm utilizing the MovieLens data set to forecast movie ratings. The refined model, incorporating movie and user effects, demonstrates a notable decrease in RMSE, making it the preferred choice for our project. This optimal model exhibits an RMSE value of **0.8641**, surpassing the initial evaluation benchmark of 0.8649 set by our project's objectives.

6 LIMITATIONS

1. RMSE could be enhanced by integrating other factors such as genre, release year etc however, we limit the scope of the project since the goal has been achieved.
2. Due to lack of access to a professional laptop and tools, automated data wrangling formulas could not be used and RMSE was instead calculated manually.